

CSE 5311: Design and Analysis of Algorithms

Red-Black Trees

Alex Dillhoff

University of Texas at Arlington

Red-Black Trees

Red-Black Trees are modified Binary Search Trees that maintain a balanced structure in order to guarantee that operations like search, insert, and delete run in $O(\log n)$ time.

Red-Black Trees

A red-black tree is a binary search tree with the following properties:

- Every node is either red or black.
- The root is black.
- Every `NULL` leaf is black.
- If a node is red, then both its children are black.
- For each node, all simple paths from the node to descendant leaves contain the same number of black nodes.

Red-Black Trees

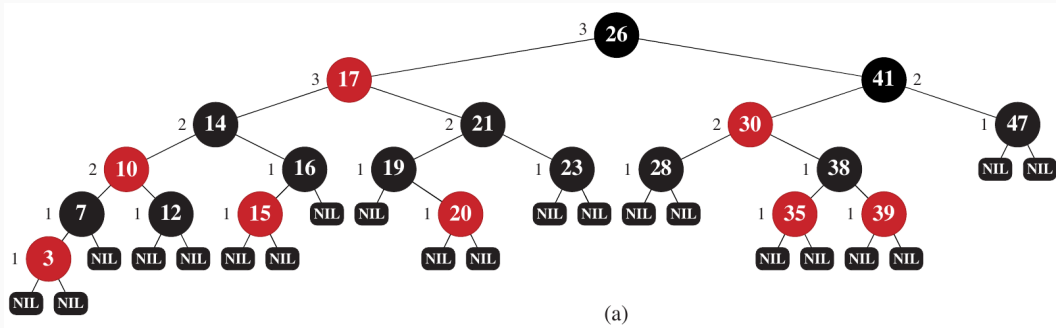


Figure 1: A red-black tree (Source: CRLS Chapter 13).

Red-Black Trees

The only structural addition we need to make over a BST is the addition of a `color` attribute to each node. This attribute can be either `RED` or `BLACK`.

Property 5 implies that the *black-height* of a tree is an important property.

This property is used to prove that the height of a red-black tree with n internal nodes is at most $2 \log(n + 1)$.

Operations

Rotate

If a Binary Search Tree is balanced, then searching for a node takes $O(\log n)$ time.

If the tree is unbalanced, then searching can take $O(n)$ time.

When items are inserted or deleted from a tree, it can become unbalanced.

Without any way to correct for this, a BST is less desirable unless the data will not change.

Rotate

When nodes are inserted or deleted into a red-black tree, the **rotation** operation is used in functions that maintain the red-black properties.

This ensures that the tree remains balanced and that operations like search, insert, and delete run in $O(\log n)$ time.

Rotate

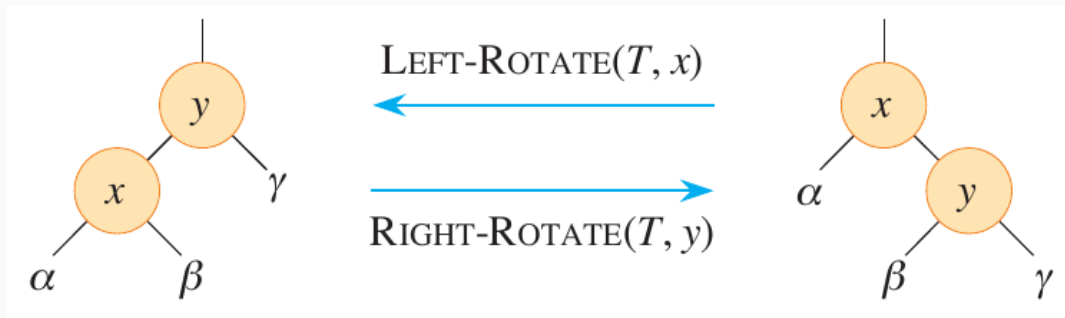


Figure 2: The rotation operation. (Source: CRLS Chapter 13).

Rotate

```
def left_rotate(self, x):  
    y = x.right  
    x.right = y.left  
    if y.left != self.nil:  
        y.left.p = x  
    y.p = x.p  
    if x.p == self.nil:  
        self.root = y  
    elif x == x.p.left:  
        x.p.left = y  
    else:  
        x.p.right = y  
    y.left = x  
    x.p = y
```

Rotate

```
def right_rotate(self, y):  
    x = y.left  
    y.left = x.right  
    if x.right != self.nil:  
        x.right.p = y  
    x.p = y.p  
    if y.p == self.nil:  
        self.root = x  
    elif y == y.p.left:  
        y.p.left = x  
    else:  
        y.p.right = x  
    x.right = y  
    y.p = x
```

Rotate




Figure 3: The rotation operation in action. (Source: CRLS Chapter 13).

Rotate

What is the upper bound on the running time of a rotation?

Rotate

What is the upper bound on the running time of a rotation?

$O(1)$. Only pointer assignments are updated.

Insert

The `insert` operation in a red-black tree starts off identically to the `insert` operation in a BST.

The new node is inserted into the tree as a leaf node.

Since the `NULL` leaf nodes must be black by definition, the added node is colored red.

Insert

Example: `insert` operation.

Insert

By adding the node and setting its color to red, we have possibly violated properties 2 and 4.

Property 2 is violated if z is the root.

Property 4 is violated if the parent of the new node is also red.

The final line of the function calls `insert_fixup` to restore the red-black properties.

Insert

Example: `insert_fixup` operation.

Insert Fixup: Case 1

Inside the **while** loop, the first and second conditions are symmetric.

One considers the case where z 's parent is a left child, and the other considers the case where z 's parent is a right child.

If z 's parent is a left child, then we start by setting y to z 's *aunt*.

Insert Fixup: Case 1

Let's investigate the first **if** statement, where y is RED.

In this case, both z 's parent and aunt are RED.

We can fix this by setting both to BLACK and setting z 's grandparent to RED.

This may violate property 2, so we set z to its grandparent and repeat the loop.

Insert Fixup: Case 1

```
if y.color == RED:  
    z.p.color = BLACK  
    y.color = BLACK  
    z.p.p.color = RED  
    z = z.p.p
```

Insert Fixup: Case 2

If y is BLACK, then we need to consider the case where z is a right child.

In this case, we set z to its parent and perform a left rotation.

This automatically results in the third case, where z is a left child.

Insert Fixup: Case 2

```
if z == z.p.right:  
    z = z.p  
    self.left_rotate(z)
```

Insert Fixup: Case 3

If z is a left child, then we set z 's parent to BLACK and its grandparent to RED.

Then we perform a right rotation on the grandparent.

Insert Fixup: Case 3

```
z.p.color = BLACK  
z.p.p.color = RED  
self.right_rotate(z.p.p)
```

Insert Fixup

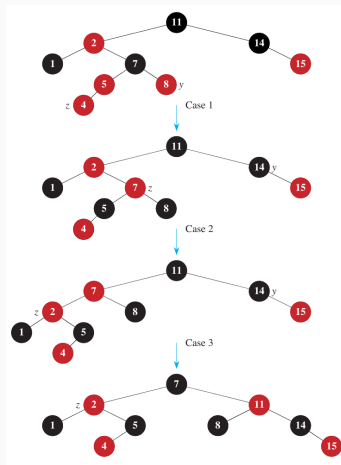


Figure 4: Insert Fixup cases. (Source: CRLS Chapter 13).

Delete

Delete

Like the `delete` operation of a BST, the `delete` operation of a RBT uses a `transplant` operation to replace the deleted node with its child.

The `transplant` operation is defined as follows.

Delete

```
def transplant(self, u, v):  
    if u.p == self.nil:  
        self.root = v  
    elif u == u.p.left:  
        u.p.left = v  
    else:  
        u.p.right = v  
    v.p = u.p
```

Delete

The full `delete` operation follows a similar structure to that of its BST counterpart.

There are a few distinct differences based on the color of the node being deleted.

Delete

```
def delete(self, z):  
    y = z  
    y_original_color = y.color
```

Delete

The first line sets y to the node to be deleted. The second line saves the color of y .

This is necessary because y will be replaced by another node, and we need to know the color of the replacement node.

Delete

The first two conditionals check if z has any children.

If there is a right child, then the z is replaced by the left child.

If there is a left child, then z is replaced by the right child.

If z has no children, then z is replaced by `NULL`.

Delete

```
if z.left == None:
    x = z.right
    self.transplant(z, z.right)
elif z.right == None:
    x = z.left
    self.transplant(z, z.left)
```

Delete

If z has two children, then we find the successor of z and set y to it.

The successor is guaranteed to have at most one child, so we can use the code above to replace y with its child.

Then we replace z with y .

Delete

```
else:
    y = self.minimum(z.right)
    y_original_color = y.color
    x = y.right
    if y != z.right: # y is farther down the tree
        self.transplant(y, y.right)
        y.right = z.right
        y.right.p = y
    else:
        x.p = y
    self.transplant(z, y)
    y.left = z.left
    y.left.p = y
    y.color = z.color
```

Delete

The procedure kept track of `y_original_color` to see if any violations occurred.

This would happen if `y` was originally `BLACK` because the `transplant` operation, or the deletion itself, could have violated the red-black properties.

If `y_original_color` is `BLACK`, then we call `delete_fixup` to restore the properties.

Delete Fixup

If the node being deleted is **BLACK**, then the following scenarios can occur.

If y is the root and a **RED** child of y becomes the new root, property 2 is violated.

Let x be a **RED** child of y , if a new parent of x is **RED**, then property 4 is violated.

Removing y may have caused a violation of property 5, since any path containing y has 1 less **BLACK** node in it.

Delete Fixup

Correcting violation 5 can be done by *transferring* the BLACK property from y to x , the node that moves into y 's original position.

This requires us to allow nodes to take on multiple counts of colors.

Delete Fixup

If x was already BLACK, it becomes double BLACK.

If it was RED, it becomes RED-AND-BLACK.

There is a good reason to this extension, as it will help us decide which case of `delete_fixup` to use.

Delete Fixup

The `delete_fixup` function will restore violations of properties 1, 2, and 4.

It is called after the `delete` operation, and it takes a single argument, `x`, which is the node that replaced the deleted node.

It performs a series of rotations and color changes to restore the violated properties.

Delete Fixup

Let's look at the `delete_fixup` function from the ground up.

It is a little more complex than `insert_fixup` because it has to handle the case where the node being deleted is **BLACK**.

In total, there are 4 distinct cases per side.

Delete Fixup

Like `insert_fixup`, it is enough to understand the first half, as the second is symmetric.

The function begins as follows, where `x` is a left child.

Delete Fixup - Case 1

```
def delete_fixup(self, x):  
    while x != self.root and x.color == BLACK:  
        if x == x.p.left:  
            w = x.p.right  
            if w.color == RED:  
                w.color = BLACK  
                x.p.color = RED  
                self.left_rotate(x.p)  
            w = x.p.right
```

Delete Fixup - Case 1

In the first case, x 's sibling w is RED.

If this is true, then w must have two BLACK subnodes.

The colors of w and x 's parent are then switched, and a left rotation is performed on x 's parent.

The result of case 1 converts to one of cases 2, 3, or 4.

Delete Fixup - Case 1




Figure 5: Delete Fixup Case 1 (Source: CRLS Chapter 13).

Delete Fixup - Case 2

```
if w.left.color == BLACK and w.right.color == BLACK:  
    w.color = RED  
    x = x.p
```

Delete Fixup - Case 2

If both of w 's subnodes are **BLACK** and both w and x are also black (actually, x is doubly **BLACK**)...

then there is an extra **BLACK** node on the path from w to the leaves.

Delete Fixup - Case 2

The colors of both x and w are switched, which leaves x with a single BLACK count and w as RED.

The extra BLACK property is transferred to x 's parent.

Delete Fixup - Case 2




Figure 6: Delete Fixup Case 2 (Source: CRLS Chapter 13).

Delete Fixup - Case 3

```
else:
    if w.right.color == BLACK:
        w.left.color = BLACK
        w.color = RED
        self.right_rotate(w)
    w = x.p.right
```

Delete Fixup - Case 3

If w is BLACK, its left child is RED, and its right child is BLACK, then the colors of w and its left child are switched.

Then a right rotation is performed on w .

This rotation moves the BLACK node to w 's position, which is now the new sibling of x .

This leads directly to case 4.

Delete Fixup - Case 3




Figure 7: Delete Fixup Case 3 (Source: CRLS Chapter 13).

Delete Fixup - Case 4

```
w.color = x.p.color  
x.p.color = BLACK  
w.right.color = BLACK  
self.left_rotate(x.p)  
x = self.root
```

Delete Fixup - Case 4

At this point, w is BLACK and its right child is RED.

Also remember that x still holds an extra BLACK count.

This last case performs color changes and a left rotation which remedy the extra BLACK count.

Delete Fixup - Case 4




Figure 8: Delete Fixup Case 4 (Source: CRLS Chapter 13).

Delete Runtime

The `delete` operation takes $O(\log n)$ time since it performs a constant number of rotations.

The `delete_fixup` operation also takes $O(\log n)$ time since it performs a constant number of color changes and at most 3 rotations.

Case 2 of `delete_fixup` could move the violation up the tree, but this would happen no more than $O(\log n)$ times.

In total, the `delete` operation takes $O(\log n)$ time.