

# CSE 4373/5373 - General Purpose GPU Programming

## GPU Pattern: Merge

---

Alex Dillhoff

University of Texas at Arlington

# Merge

- The **merge** operation takes two sorted subarrays and combines them into a single sorted array.
- You may be familiar with this approach from studying Divide and Conquer algorithms.
- Parallelizing the merge operation is a non-trivial task and will require the use of a few new techniques.

# Merge

In the context of GPU programming, we will learn about the following techniques:

- Dynamic input data identification
- Data locality
- Buffer management schemes

# Merge

- Using merge from "Perfectly load-balanced, optimal, stable, parallel merge" (Siebert et al., 2013).
- Compute which values are needed in each merge step,
- Use a parallel kernel to compute the merge.
- These steps can be computed by each thread independently.

# Co-rank Function

The key to this implementation is the **co-ranking function**.

- Compute the range of indices needed from two input values to produce a given output value.
- Avoids the need to merge the two input arrays explicitly

# Co-rank Function

When merging two sorted arrays, we can observe that the output index  $0 \leq k < m + n$  comes from either

1.  $0 \leq i < m$  from input A or
2.  $0 \leq j < n$  from input B.

# Co-rank Function

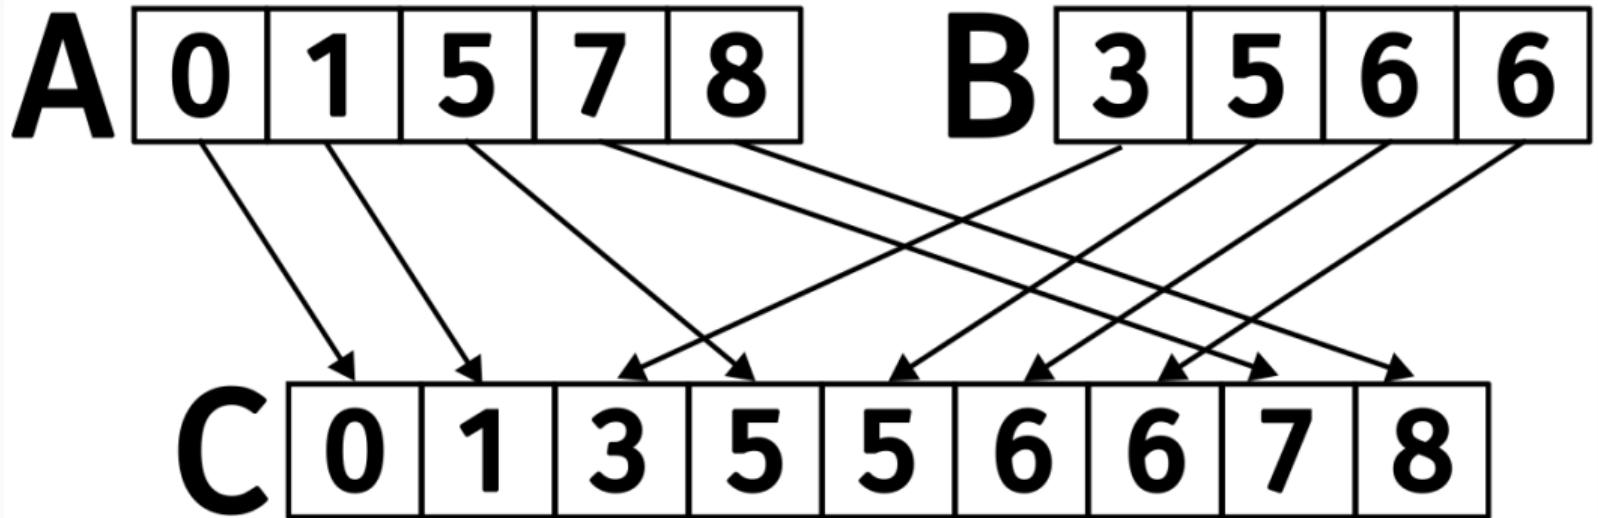


Figure 1: Merging two sorted arrays.

# Co-rank Function

The element at  $k = 3$  comes from  $A[2]$ , so  $i = 2$ .

It must be that  $k = 3$  is the result of merging the first  $i = 2$  elements of  $A$  with the first  $j = k - i$  elements of  $B$ .

# Co-rank Function

This works both ways: for  $k = 6$ , the value is taken from  $B[3]$ , so  $j = 3$ , and the result is the merge of the first  $i = k - j$  elements of  $A$  with the first  $j = 3$  elements of  $B$ .

# Co-rank Function

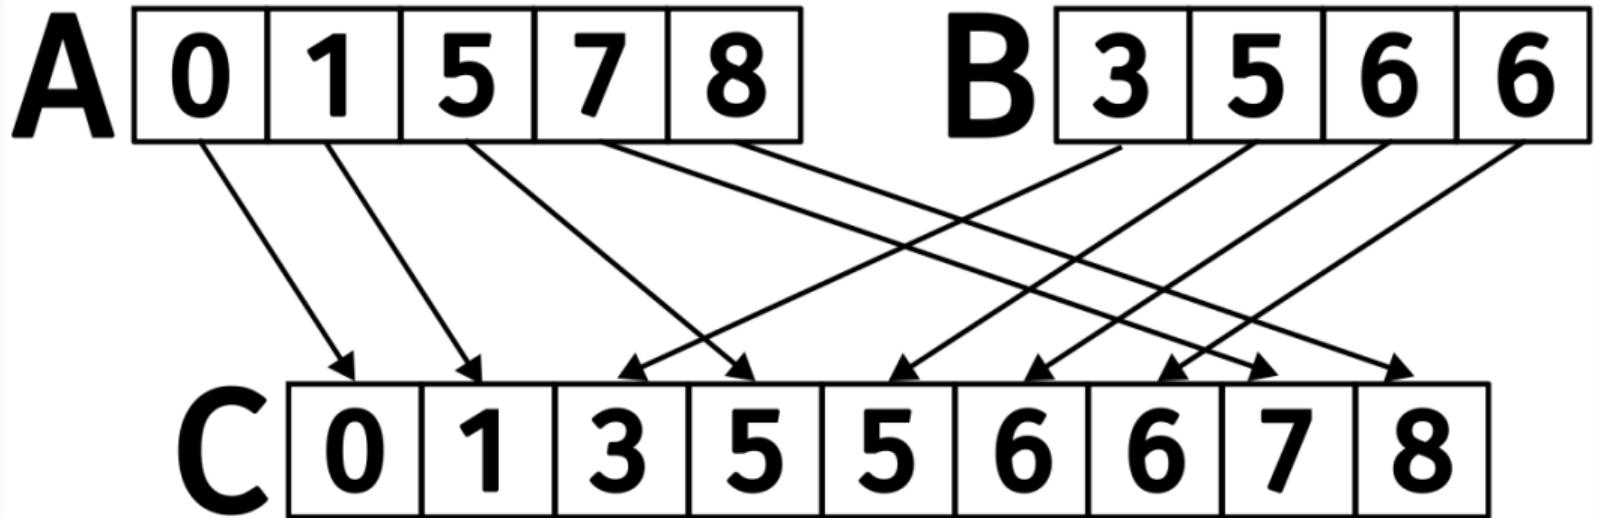


Figure 2: Merging two sorted arrays.

# Co-rank Function

## Lemma 1

For any  $k$ ,  $0 \leq k < m + n$ , there exists a unique  $i$ ,  $0 \leq i \leq m$ , and a unique  $j$ ,  $0 \leq j \leq n$ , with  $i + j = k$  such that

1.  $i = 0$  or  $A[i - 1] \leq B[j]$  and
2.  $j = 0$  or  $B[j - 1] < A[i]$ .

# Co-rank Function

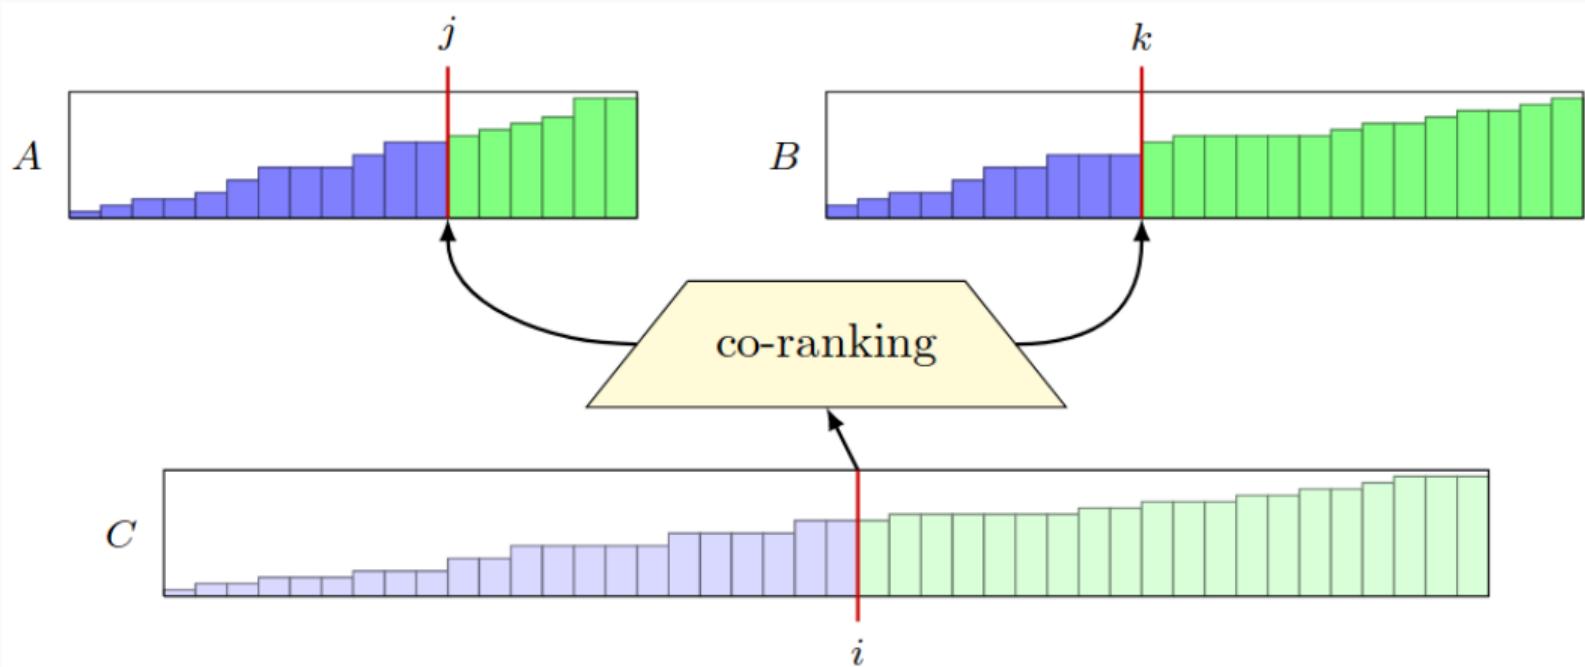


Figure 3: Visualization of the co-rank function.

# Implementation

---

# Implementation of Co-rank

Given the rank  $k$  of an element in an output array  $C$  and two input arrays  $A$  and  $B$ ...

the co-rank function  $f$  returns the co-rank value for the corresponding element in  $A$  and  $B$ .

# Implementation of Co-rank

How would the co-rank function be used in the example above?

# Implementation of Co-rank

How would the co-rank function be used in the example above?

Given two threads, let thread 1 compute the co-rank for  $k = 4$ .

This would return  $i = 3$  and  $j = 1$ .

# Implementation of Co-rank

We can verify this result using Lemma 1.

$$A[2] = 5 \leq B[1] = 5 \text{ and } B[0] = 3 < A[3] = 7.$$

# Implementation of Co-rank

## Initialization

```
int co_rank(int k, int *A, int m, int *B, int n) {  
    int i = min(k, m);  
    int j = k - i;  
    int i_low = max(0, k-n);  
    int j_low = max(0, k-m);  
    int delta;  
    bool active = true;
```

# Implementation of Co-rank

```
while (active) {  
    if (i > 0 && j < n && A[i-1] > B[j]) {  
        delta = (i - i_low + 1) / 2;  
        j_low = j;  
        i -= delta;  
        j += delta;  
    } else if (j > 0 && i < m && B[j-1] >= A[i]) {  
        delta = (j - j_low + 1) / 2;  
        i_low = i;  
        j -= delta;  
        i += delta;  
    } else {  
        active = false;  
    }  
}
```

# Co-rank Example

Consider running a merge kernel across 3 threads where each thread takes 3 sequential output values.

Use the co-rank function to compute the co-rank values for  $k = 3$  and  $k = 6$ , simulating the tasks for the second and third threads.

# Co-rank Example

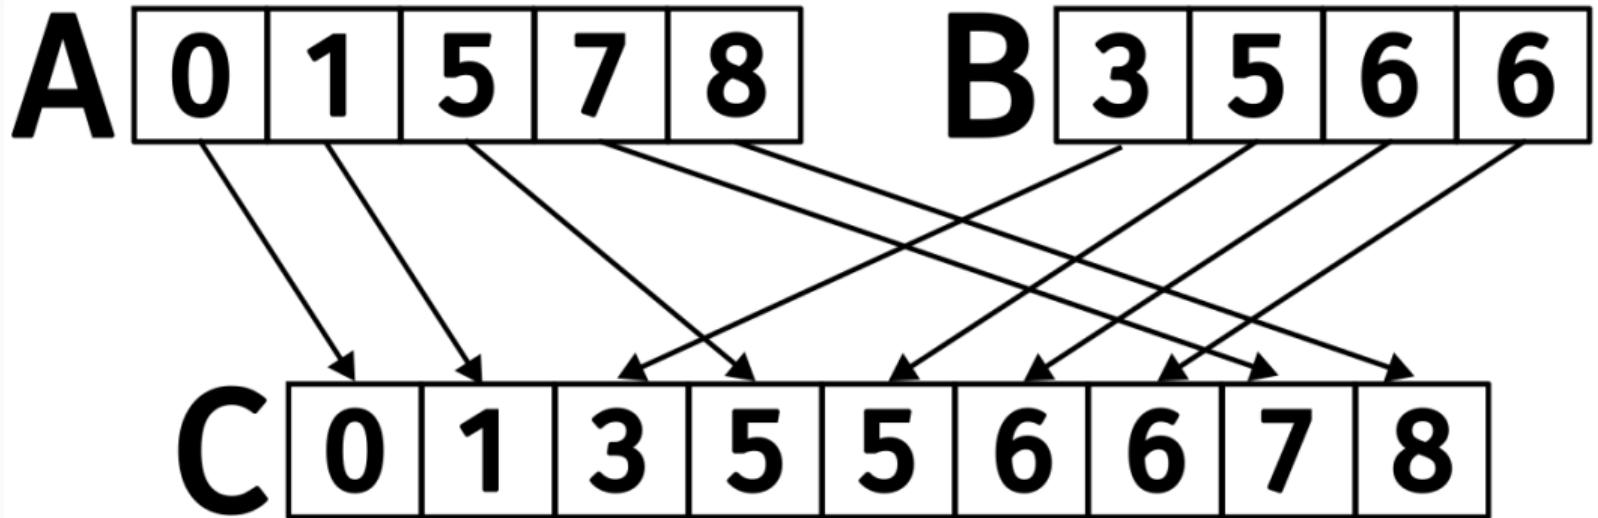


Figure 4: Merging two sorted arrays.

## Co-rank Example

The values for  $k = 3$  should be  $i = 2$  and  $j = 1$ , for reference.

All values below these indices would be used by the first thread.

# Parallel Kernel

We can now implement a basic parallel merge kernel.

- Each thread is responsible for determining how many elements it will be responsible for merging.
- The range of input values is determined via two calls to `co_rank`, one for the starting and ending point.

# Parallel Kernel

```
__global__ void merge_basic_kernel(int *A, int m, int *B, int n, int *C) {
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    int elementsPerThread = ceil((m + n) / (blockDim.x * gridDim.x));
    int k_curr = tid * elementsPerThread; // start output index
    int k_next = min((tid + 1) * elementsPerThread, m + n); // end output index
    int i_curr = co_rank(k_curr, A, m, B, n);
    int i_next = co_rank(k_next, A, m, B, n);
    int j_curr = k_curr - i_curr;
    int j_next = k_next - i_next;
    merge_sequential(&A[i_curr], i_next - i_curr,
                      &B[j_curr], j_next - j_curr, &C[k_curr]);
}
```

# Parallel Kernel

Bottlenecks?

# Parallel Kernel

## Bottlenecks?

1. Memory accesses to input arrays are not coalesced.
2. Binary search in the co-rank function is not coalesced.

# Tiled Merge

---

# Tiled Merge

- Improve memory accesses by having the threads transfer data from global memory to shared memory in a coalesced manner.
- The higher latency operation will be coalesced.
- The data in shared memory may be accessed out of order, but the latency is much lower.

# Tiled Merge

- The subarrays from  $A$  and  $B$  that are used by adjacent threads are also adjacent in memory.
- By considering block-level subarrays, we can ensure that the data is coalesced.
- This is the idea behind the tiled merge algorithm.

# Tiled Merge

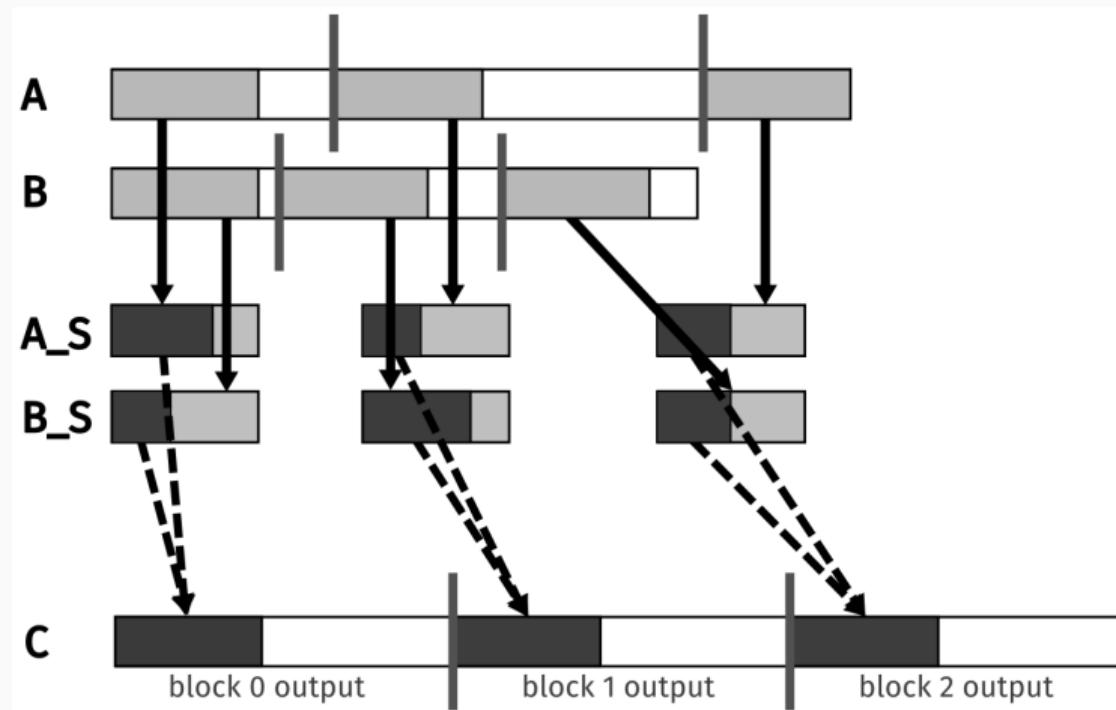


Figure 5: Visualization of the tiled merge algorithm.

# Tiled Merge

- The shared memory blocks cannot store the entire range of data needed.

# Tiled Merge

- The shared memory blocks cannot store the entire range of data needed.
- In each iteration, the threads in a block will load a new set of data from global memory to shared memory.

# Tiled Merge

- The shared memory blocks cannot store the entire range of data needed.
- In each iteration, the threads in a block will load a new set of data from global memory to shared memory.
- If they collectively load  $2n$  elements, only  $n$  elements will be used in the merge operation.

# Tiled Merge

- The shared memory blocks cannot store the entire range of data needed.
- In each iteration, the threads in a block will load a new set of data from global memory to shared memory.
- If they collectively load  $2n$  elements, only  $n$  elements will be used in the merge operation.
- This is because in the worst case, all elements going to the output array will come from one of the two input arrays.

# Tiled Merge

Each block will use a portion of both  $A_s$  and  $B_s$  to compute the merge.

This is shown with dotted lines going from the shared memory to the output array.

# Tiled Merge Kernel: Part 1

---

# Tiled Merge: Part 1

Establish shared memroy and output boundaries.

```
--global__ void merge_tiled_kernel(int *A, int m, int n,
                                    int *C, int tile_size) {
    extern __shared__ int shareAB[];
    int *A_s = &shareAB[0];
    int *B_s = &shareAB[tile_size];
    int C_curr = blockIdx.x * ceil((m+n)/gridDim.x);
    int C_next = min((blockIdx.x+1) * ceil((m+n)/gridDim.x), m+n);
```

# Tiled Merge: Part 1

Only the first thread needs to compute the co-rank values.

```
if (threadIdx.x == 0) {  
    // Block-level co-rank values will be  
    // available to all threads in the block  
    A_s[0] = co_rank(C_curr, A, m, B, n);  
    A_s[1] = co_rank(C_next, A, m, B, n);  
}  
__syncthreads();
```

# Tiled Merge: Part 1

Compute  $j$  values based on  $i$  values.

```
int A_curr = A_s[0];
int A_next = A_s[1];
int B_curr = C_curr - A_curr;
int B_next = C_next - A_next;
__syncthreads();
```

# Tiled Merge Kernel: Part 2

---

## Tiled Merge: Part 2

The second part of the kernel is responsible for loading the input data into shared memory.

This is done in a coalesced manner, as the threads in a block will load a contiguous section of the input arrays.

## Tiled Merge: Part 2

Compute lengths of each section and number of iterations for coarsening.

```
int counter = 0;  
int C_length = C_next - C_curr;  
int A_length = A_next - A_curr;  
int B_length = B_next - B_curr;  
int total_iteration = ceil(C_length / tile_size);  
int C_completed = 0;  
int A_consumed = 0;  
int B_consumed = 0;
```

# Tiled Merge: Part 2

Copy data from global memory to shared memory.

```
while (counter < total_iteration) {
    for (int i = 0; i < tile_size; i += blockDim.x) {
        if (i + threadIdx.x < A_length - A_consumed) {
            A_s[i + threadIdx.x] = A[A_curr + A_consumed + i + threadIdx.x];
        }
        if (i + threadIdx.x < B_length - B_consumed) {
            B_s[i + threadIdx.x] = B[B_curr + B_consumed + i + threadIdx.x];
        }
    }
    __syncthreads();
}
```

# Tiled Merge Kernel: Part 3

---

## Tiled Merge: Part 3

With the input in shared memory, each thread will divide up this input and merge their respective sections in parallel.

This is done by calculating the `c_curr` and `c_next` first, which is the output section of the thread.

Using those boundaries, two calls to `co_rank` will determine the input sections the thread.

# Tiled Merge: Part 3

Compute output section for the current thread.

```
int c_curr = threadIdx.x * (tile_size / blockDim.x);
int c_next = (threadIdx.x + 1) * (tile_size / blockDim.x);
c_curr = (c_curr <= C_length - C_completed) ?
          c_curr : C_length - C_completed;
c_next = (c_next <= C_length - C_completed) ?
          c_next : C_length - C_completed;
```

# Tiled Merge: Part 3

Compute co-rank values for the current section.

```
int a_curr = co_rank(c_curr,
                      A_s, min(tile_size, A_length - A_consumed),
                      B_s, min(tile_size, B_length - B_consumed));
int b_curr = c_curr - a_curr;
int a_next = co_rank(c_next,
                      A_s, min(tile_size, A_length - A_consumed),
                      B_s, min(tile_size, B_length - B_consumed));
int b_next = c_next - a_next;
```

# Tiled Merge: Part 3

Merge the data and compute how much data was used.

```
    merge_sequential(&A_s[a_curr], a_next - a_curr,  
                      &B_s[b_curr], b_next - b_curr,  
                      &C[C_urr + C_completed + c_curr]);  
  
    counter++;  
  
    C_completed += tile_size;  
  
    A_consumed += co_rank(tile_size A_s, tile_size, B_s, tile_size);  
    B_consumed = C_completed - A_consumed;  
    __syncthreads();  
}  
}
```

# Kernel Walkthrough

We have two input arrays  $A = [1, 3, 5, 7, 9]$  and  $B = [2, 4, 6, 8, 10]$ .

- The output array C will have 10 elements.
- Use 2 blocks and 4 threads per block.
- The tile size is 4.
- With 10 elements and 2 blocks, each block is responsible for 5 elements.

# Kernel Walkthrough

The main **while** loop will need to iterate twice to cover the entire output array.

- The first iteration will load the first 4 elements of A and B into shared memory.

# Kernel Walkthrough

The main **while** loop will need to iterate twice to cover the entire output array.

- The first iteration will load the first 4 elements of A and B into shared memory.
- Next, each thread divides the input tiles by running the co-rank function on the data that is in shared memory.
- The computed indices are the boundaries between each thread.

# Kernel Walkthrough

- In each iteration, a block is responsible for 4 elements.
- Given that we have 4 threads per block, each thread will be responsible for 1 output element per iteration.

# Kernel Walkthrough

- In each iteration, a block is responsible for 4 elements.
- Given that we have 4 threads per block, each thread will be responsible for 1 output element per iteration.
- **Thread 0:**  $c\_curr = 0$  and  $c\_next = 2$ .
- This results in  $a\_curr = 0$ ,  $b\_curr = 0$ ,  $a\_next = 1$ , and  $b\_next = 1$ .
- The merge operation will then be performed on the first element of A and B.

# Kernel Analysis

---

# Analysis

The tiled kernel is an improvement due to...

1. coalescing of global memory accesses
2. Shared memory use for reduced latency

# Analysis

What is the bottleneck of this kernel?

# Analysis

What is the bottleneck of this kernel?

Only half the data loaded into shared memory is used, leading to wasted memory bandwidth.

# Analysis

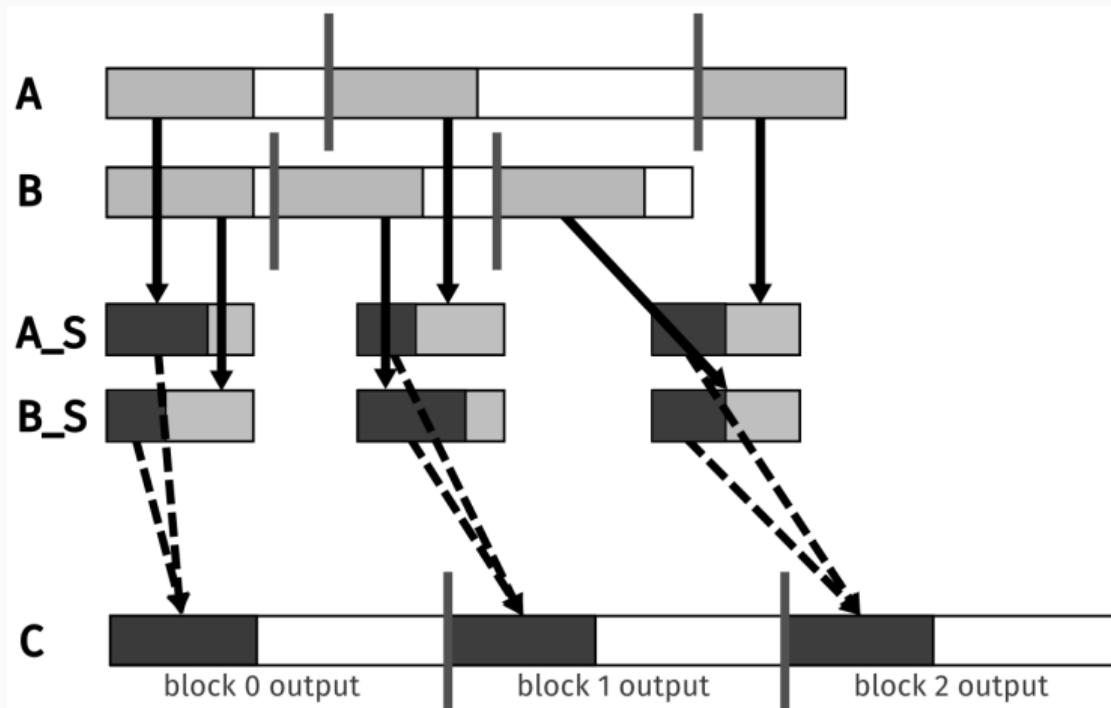


Figure 6: Visualization of the tiled merge algorithm.

# Circular Buffers

---

# Circular Buffers

The tiled version leaves half of the data unused every iteration.

A **circular buffer** can be implemented to reuse the data that was loaded into shared memory.

# Circular Buffers

Instead of writing over the shared memory values each iteration, the data stays in memory.

A portion of new data is loaded into shared memory, reducing the number of reads from global memory.

# Circular Buffers

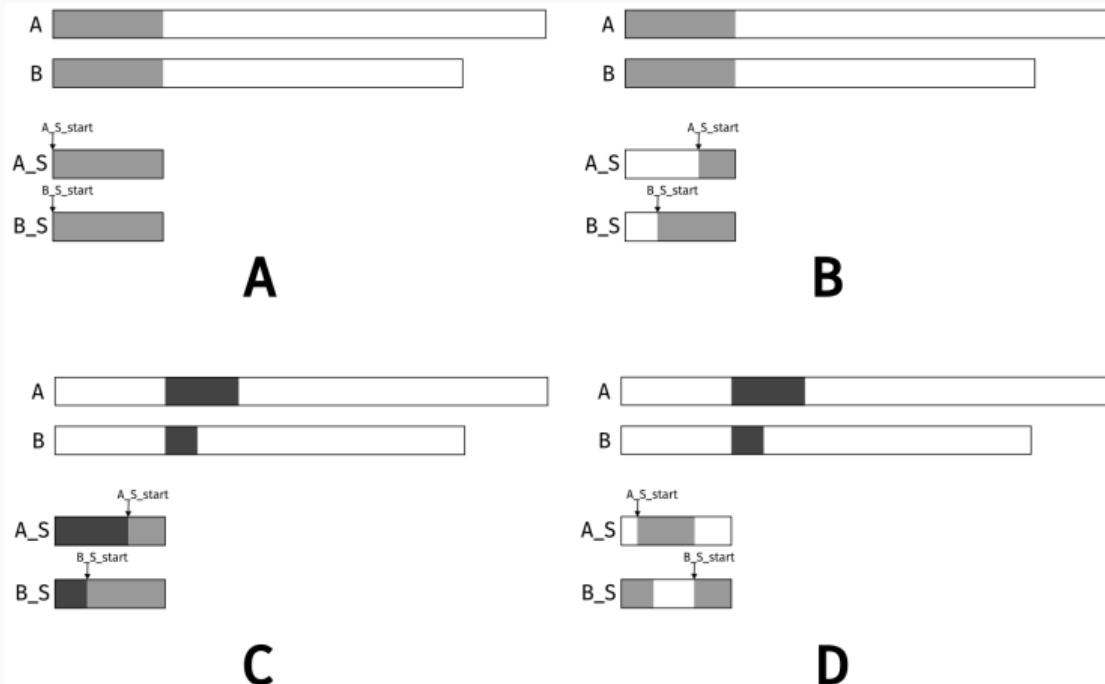


Figure 7: Visualization of the circular buffer.

# Circular Buffers

- Part A shows the initial layout.
- Part B shows a blank portion that depicts what was used, the light gray represents what remains.
- Part C shows the new data loaded into shared memory.
- Part D shows the next iteration.

# Implementation

- `A_consumed` is used to keep track of how many new elements need to be read.
- The `co_rank` and `merge_sequential` functions need to be updated to work with circular buffers.
- It is easier to treat the shared memory as an extended array...
- this avoids situations where the `next` index is less than the `current` index.

# Updating the Co-rank Function

```
int co_rank_circular(int k, int *A, int m,
                      int *B, int n,
                      int A_S_start, int B_S_start,
                      int tile_length) {
    int i = min(k, m);
    int j = k - i;
    int i_low = max(0, k-n);
    int j_low = max(0, k-m);
    int delta;
    bool active = true;
```

# Updating the Co-rank Function

```
while (active) {  
    int i_cir = (A_S_start + i) % tile_length;  
    int j_cir = (B_S_start + j) % tile_length;  
    int i_m_1_cir = (A_S_start + i - 1) % tile_length;  
    int j_m_1_cir = (B_S_start + j - 1) % tile_length;
```

# Updating the Co-rank Function

```
if (i > 0 && j < n && A[i_m_1_cir] > B[j_cir]) {  
    delta = ((i - i_low + 1) >> 1);  
    j_low = j;  
    i -= delta;  
    j += delta;  
} else if (j > 0 && i < m && B[j_m_1_cir] >= A[i_cir]) {  
    delta = ((j - j_low + 1) >> 1);  
    i_low = i;  
    j -= delta;  
    i += delta;  
} else {  
    active = false;  
}  
}  
  
return i;
```

# Updating the Co-rank Function

In this updated version of the co-rank function...

the user only needs to provide the start indices for the shared memory arrays along with the tile size.

# Updating the Merge Function

```
void merge_sequential_circular(int *A, int m,
                                int *B, int n,
                                int *C, int A_S_start,
                                int B_S_start, int tile_size) {

    int i = 0;
    int j = 0;
    int k = 0;
```

# Updating the Merge Function

```
while (i < m && j < n) {  
    int i_cir = (A_S_start + i) % tile_size;  
    int j_cir = (B_S_start + j) % tile_size;  
    if (A[i_cir] <= B[j_cir]) {  
        C[k] = A[i_cir];  
        i++;  
    } else {  
        C[k] = B[j_cir];  
        j++;  
    }  
    k++;  
}
```

# Updating the Merge Function

```
if (i == m) {  
    while (j < n) {  
        int j_cir = (B_S_start + j) % tile_size;  
        C[k] = B[j_cir];  
        j++;  
        k++;  
    }  
} else {  
    while (i < m) {  
        int i_cir = (A_S_start + i) % tile_size;  
        C[k] = A[i_cir];  
        i++;  
        k++;  
    }  
}
```

# Updating the Merge Function

This too acts as a drop-in replacement, only requiring the start indices of the shared memory arrays and the tile size.

# Circular Buffer Kernel

```
int c_curr = threadIdx.x * (tile_size / blockDim.x);
int c_next = (threadIdx.x + 1) * (tile_size / blockDim.x);
c_curr = (c_curr <= C_length - C_completed) ?
          c_curr : C_length - C_completed;
c_next = (c_next <= C_length - C_completed) ?
          c_next : C_length - C_completed;
```

# Circular Buffer Kernel

```
int a_curr = co_rank_circular(c_curr,
                               A_s, min(tile_size, A_length - A_consumed),
                               B_s, min(tile_size, B_length - B_consumed),
                               A_curr, B_curr, tile_size);

int b_curr = c_curr - a_curr;
int a_next = co_rank_circular(c_curr,
                               A_s, min(tile_size, A_length - A_consumed),
                               B_s, min(tile_size, B_length - B_consumed),
                               A_curr, B_curr, tile_size);

int b_next = c_next - a_next;
```

# Circular Buffer Kernel

```
merge_sequential_circular(A_s, a_next - a_curr,  
                           B_s, b_next - b_curr,  
                           &C[C_urr + C_completed + c_curr],  
                           A_S_start + A_curr, B_S_start + B_curr, tile_size);
```

# Circular Buffer Kernel

```
// Compute the indices that were used
counter++;
A_S_consumed = co_rank_circular(min(tile_size, C_length - C_completed),
                                  A_s, min(tile_size, A_length - A_consumed),
                                  B_s, min(tile_size, B_length - B_consumed),
                                  A_S_start, B_S_start, tile_size);
B_S_consumed = min(tile_size, C_length - C_completed) - A_S_consumed;
A_consumed += A_S_consumed;
C_completed += min(tile_size, C_length - C_completed);
B_consumed = C_completed - A_consumed;
```

# Circular Buffer Kernel

```
// Update the start indices for the next iteration
A_S_start = (A_S_start + A_S_consumed) % tile_size;
B_S_start = (B_S_start + B_S_consumed) % tile_size;
__syncthreads();
```

# Coarsening

---

# Thread Coarsening

- The kernels presented already utilize thread coarsening.
- Each thread is responsible for a range of output values.
- Each thread was only responsible for a single output value.